

Grammar-Based Concept Alignment for Domain-Specific Machine Translation

08.09.2021

Arianna Masciolini and Aarne Ranta

Context

- ✚ In **domain-specific MT**, precision is often more important than coverage

Context

- ✦ In **domain-specific MT**, precision is often more important than coverage
- ✦ grammar-based pipelines (cf. GF) provide strong guarantees of **grammatical correctness**

Context

- ❖ In **domain-specific MT**, precision is often more important than coverage
- ❖ grammar-based pipelines (cf. GF) provide strong guarantees of **grammatical correctness**
- ❖ **lexical exactness** is as important as grammaticality
 - ❖ need for high-quality **translation lexica** preserving semantics *and* morphological correctness

Translation lexica

- ❖ Often built **manually**
 - ❖ **time** consuming
 - ❖ significant **linguistic knowledge** required

Translation lexica

- ❖ Often built **manually**
 - ❖ **time** consuming
 - ❖ significant **linguistic knowledge** required
- ❖ need for at least partial **automation**
 - ❖ **example parallel data** required

A parallel corpus

Alice thought she might as well wait, as she had nothing else to do, and perhaps after all it might tell her something worth hearing.

For some minutes it puffed away without speaking, but at last it unfolded its arms, took the hookah out of its mouth again, and said, 'So you think you're changed, do you?'

'I'm afraid I am, sir,' said Alice; 'I can't remember things as I used--and I don't keep the same size for ten minutes together!'

Alice pensò che poteva aspettare, perchè non aveva niente di meglio da fare, e perchè forse il Bruco avrebbe potuto dirle qualche cosa d'importante.

Per qualche istante il Bruco fumò in silenzio, finalmente sciolse le braccia, si tolse la pipa di bocca e disse: — E così, tu credi di essere cambiata?

— Ho paura di sì, signore, — rispose Alice. — Non posso ricordarmi le cose bene come una volta, e non rimango della stessa statura neppure per lo spazio di dieci minuti!

From Lewis Carroll, *Alice's Adventures in Wonderland*. Parallel text at paralleltext.io

Types of alignment

Word alignment:

Alice thought she might as well wait, as she had **nothing** else to do, and perhaps after all it might tell her something worth hearing.

Alice pensò che poteva aspettare, perchè non aveva **niente** di meglio da fare, e perchè forse il Bruco avrebbe potuto dirle qualche cosa d'importante.

Phrase alignment:

Alice thought she might as well wait, as she had **nothing else to do**, and perhaps after all it might tell her something worth hearing.

Alice pensò che poteva aspettare, perchè non aveva **niente di meglio da fare**, e perchè forse il Bruco avrebbe potuto dirle qualche cosa d'importante.

Approaches to automation

statistical (e.g. IBM models)

require **large amounts of data**

works with **raw** data

correspondences between **strings**

“fixed” level of abstraction
(**word** or **phrase**)

syntax-based

work consistently well even on
individual sentence pairs

requires the data to be **analyzed**

correspondences between
grammatical objects

all levels of abstraction →
concept alignment

Our approach

- ❖ Inconsistencies between different grammar formalisms → translation lexicon implemented in **GF**

Our approach


- ❖ Inconsistencies between different grammar formalisms → translation lexicon implemented in **GF**
- ❖ lack of robust constituency parsers while high-quality analysis is crucial → **UD** parsing (UDPipe)

Our approach


- ❖ Inconsistencies between different grammar formalisms → translation lexicon implemented in **GF**
- ❖ lack of robust constituency parsers while high-quality analysis is crucial → **UD** parsing (UDPipe)
- ❖ gf-ud for conversion

Our approach

- ❖ Inconsistencies between different grammar formalisms → translation lexicon implemented in **GF**
- ❖ lack of robust constituency parsers while high-quality analysis is crucial → **UD** parsing (UDPipe)
- ❖ **gf-ud** for conversion



Universal Dependencies



text = she missed the boat

1 she she PRON _ _ 2 nsubj _ _

2 missed miss VERB _ _ 0 root _ _

3 the the DET _ _ 4 det _ _

4 boat boat NOUN _ _ 2 obj _

2 missed miss VERB _ _ 0 root _ _

1 she she PRON _ _ 2 nsubj _ _


4 boat boat NOUN _ _ 2 obj _

3 the the DET _ _ 4 det _ _

Graphical, CoNNL-U and Rose Tree representation of the same UD tree.

- Framework for cross-linguistically consistent grammatical annotation

Universal Dependencies



text = she missed the boat

1 she she PRON _ _ 2 nsubj _ _

2 missed miss VERB _ _ 0 root _ _

3 the the DET _ _ 4 det _ _

4 boat boat NOUN _ _ 2 obj _

2 missed miss VERB _ _ 0 root _ _

1 she she PRON _ _ 2 nsubj _ _


4 boat boat NOUN _ _ 2 obj _

3 the the DET _ _ 4 det _ _

Graphical, CoNLL-U and Rose Tree representation of the same UD tree.

- ❑ Framework for cross-linguistically consistent grammatical annotation
- ❑ dependency-labelled links between words (head-dependent pairs)

Universal Dependencies



text = she missed the boat

1 she she PRON _ _ 2 nsubj _ _

2 missed miss VERB _ _ 0 root _ _

3 the the DET _ _ 4 det _ _

4 boat boat NOUN _ _ 2 obj _

2 missed miss VERB _ _ 0 root _ _

1 she she PRON _ _ 2 nsubj _ _


4 boat boat NOUN _ _ 2 obj _

3 the the DET _ _ 4 det _ _

Graphical, CoNLL-U and Rose Tree representation of the same UD tree.

- ❑ Framework for cross-linguistically consistent grammatical annotation
- ❑ dependency-labelled links between words (head-dependent pairs)
- ❑ cannot be used for target language generation

Grammatical Framework



- ❑ Constituency grammar formalism for **multilingual grammars** (one abstract syntax + a concrete syntax per language)
- ❑ compilation-like translation (parsing + linearization)

Concept Extraction

Definitions

Concept: semantic unit of compositional translation expressed by a word or construction, conceived as a lemma equipped with morphological variations.

Definitions



Concept: semantic unit of compositional translation expressed by a word or construction, conceived as a lemma equipped with morphological variations.

Alignment: tuple of equivalent concrete expressions in different languages; represents a concept.

Extraction algorithm



```
procedure EXTRACT(criteria, (t, u))  
    alignments =  $\emptyset$   
    if (t, u) matches any alignment criteria then  
        alignments += (t, u)  
        for (t', u') in SORT(SUBTS(t))  $\times$  SORT(SUBTS(u))  
    do  
        extract(criteria, (t', u'))  
    return alignments
```

Matching UD labels



- ❑ $\langle \textit{she missed the boat, ha perso il treno} \rangle$
- ❑ $\langle \textit{missed the boat, perso il treno} \rangle$
- ❑ $*\langle \textit{the boat, il treno} \rangle$
- ❑ $\langle \textit{the, il} \rangle$

Matching UD labels





- ❑ $\langle \textit{she missed the boat, ha perso il treno} \rangle$
- ❑ $\langle \textit{missed the boat, perso il treno} \rangle$
- ❑ $*\langle \textit{the boat, il treno} \rangle$
- ❑ $\langle \textit{the, il} \rangle$

Simple improvement: aligning heads of matching subtrees



- ❑ $\langle \textit{she missed the boat, ha perso il treno} \rangle, \langle \textit{missed the boat, perso il treno} \rangle \rightarrow \langle \textit{missed, ha perso} \rangle$ (including the auxiliary)
- ❑ $\langle \textit{the boat, il treno} \rangle \rightarrow * \langle \textit{boat, treno} \rangle$

POS equivalence





❖ more reliable **ignoring function words**

POS equivalence





- ❖ more reliable **ignoring function words**
- ❖ in this case, basically same results as when matching labels

POS equivalence



- ❑ more reliable **ignoring function words**
- ❑ in this case, basically same results as when matching labels
- ❑ can increase recall when labels do not coincide

POS equivalence



- ❑ more reliable **ignoring function words**
- ❑ in this case, basically same results as when matching labels
- ❑ can increase recall when labels do not coincide
- ❑ can increase precision if used **in conjunction with labels**

Known translation divergence

Divergence: systematic cross-linguistic distinction.

Known translation divergence

Divergence: systematic cross-linguistic distinction.

❖ categorial

❖ ⟨*Gioara listens **distractedly**, Gioara lyssnar **distraherad***⟩

❖ ⟨*Herbert completed his **doctoral** thesis, Herbert ha completato la sua tesi **di dottorato***⟩

❖ conflational

❖ ⟨*Filippo is interested in **game development**, Filippo är intresserad av **spelutveckling***⟩

❖ structural

❖ ⟨*I called **Francesco**, Ho telefonato **a Francesco***⟩

❖ head swapping

❖ ⟨*Anna **usually** goes for walks, Anna **brukar** promenera*⟩

❖ thematic

❖ ⟨***Yana** likes **books**, **A Yana** piacciono **i libri***⟩

Known alignment

- ▣ Allows using CA in conjunction with statistical tools

Known alignment

- ▣ Allows using CA in conjunction with statistical tools
- ▣ iterative application

Searching for specific patterns

- ❑ gf-ud pattern matching to look for specific syntactic patterns

Searching for specific patterns

- ❑ gf-ud pattern matching to look for specific syntactic patterns
- ❑ possible generalization via pattern replacement

Searching for specific patterns

- ❖ gf-ud pattern matching to look for specific syntactic patterns
- ❖ possible generalization via pattern replacement

Example predication patterns:

- ❖ $\langle \textit{she missed the boat, ha perso il treno} \rangle \rightarrow \langle [\textit{subj}] \textit{ missed} [\textit{obj}], \textit{ ha perso} [\textit{obj}] \rangle$
- ❖ $\langle \textit{she told you that, hon berättade det för dig} \rangle \rightarrow \langle [\textit{subj}] \textit{ told} [\textit{iobj}] [\textit{obj}], [\textit{subj}] \textit{ berättade} [\textit{obj}] \textit{ för} [\textit{obl}] \rangle$

Grammar rules generation

Requirements

- ✚ aligned UD trees

Requirements

- ❑ aligned UD trees
- ❑ gf-ud

Requirements

- ❑ aligned UD trees
- ❑ gf-ud
- ❑ **morphological dictionaries**

Requirements

- ❑ aligned UD trees
- ❑ gf-ud
- ❑ **morphological dictionaries**
- ❑ **extraction grammar**

Morphological dictionaries

Purely morphological unilingual dictionaries.

Example:

```
...  
lin morphologic_A =  
    mkAMost "morphologic" "morphologicly" ;  
lin morphological_A =  
    mkAMost "morphological" "morphologically" ;  
lin morphology_N =  
    mkN "morphology" "morphologies" ;  
...
```


Extraction grammar

Defines the syntactic categories and functions to build lexical entries.

Example (prepositional NPs):

PrepNP : Prep -> NP -> PP # case head

Lexical rules

Abstract:

```
fun in_the_field__inom_området_PP : PP ;
```

English concrete:

```
lin in_the_field__inom_område_PP =  
  PrepNP in_Prep (DetCN the_Det (UseN field_N))
```

Evaluation

Evaluating extraction

UD tree alignments are evaluated:

- ❖ independently from the quality of UD parsing (100-sentence subset of the manually annotated PUD corpus)
- ❖ on raw text (DMI and CSE course plans corpora)

Evaluating extraction

UD tree alignments are evaluated:

- ❖ independently from the quality of UD parsing (100-sentence subset of the manually annotated PUD corpus)
- ❖ on raw text (DMI and CSE course plans corpora)

Metrics:

- ❖ % correct alignments
- ❖ % “useful” alignments

Results on PUD corpus

	CE		fast_align (100 sentences)		fast_align (full dataset)	
	en-it	en-sv	en-it	en-sv	en-it	en-sv
distinct alignments	536	638	1242	1044	1286	1065
correct	392 (73%)	514 (80%)	346 (28%)	538 (52%)	540 (42%)	677 (64%)
usable in MT	363 (68%)	503 (79%)	316 (25%)	525 (50%)	510 (40%)	666 (63%)

Results on PUD corpus

	CE		fast_align (100 sentences)		fast_align (full dataset)	
	en-it	en-sv	en-it	en-sv	en-it	en-sv
distinct alignments	536	638	1242	1044	1286	1065
correct	392 (73%)	514 (80%)	346 (28%)	538 (52%)	540 (42%)	677 (64%)
usable in MT	363 (68%)	503 (79%)	316 (25%)	525 (50%)	510 (40%)	666 (63%)

- ❑ CE module compared with fast_align, so extracting only one-to-many and many-to-one alignments

Results on PUD corpus

	CE		fast_align (100 sentences)		fast_align (full dataset)	
	en-it	en-sv	en-it	en-sv	en-it	en-sv
distinct alignments	536	638	1242	1044	1286	1065
correct	392 (73%)	514 (80%)	346 (28%)	538 (52%)	540 (42%)	677 (64%)
usable in MT	363 (68%)	503 (79%)	316 (25%)	525 (50%)	510 (40%)	666 (63%)

- ❖ CE module compared with fast_align, so extracting only one-to-many and many-to-one alignments
- ❖ CE has much higher precision, even when fast_align is trained on full 1000-sentence corpus

Results on course plans corpora

	PUD (100 sentences)		course plans	
	en-it	en-sv	DMI (881 sentences)	CSE (539 sentences)
distinct alignments	1197	1325	1823	1950
correct	916 (77%)	1112 (85%)	1205 (66%)	1269 (66%)
usable in MT	880 (74%)	1099 (84%)	1157 (63%)	1248 (64%)

Results on course plans corpora

	PUD (100 sentences)		course plans	
	en-it	en-sv	DMI (881 sentences)	CSE (539 sentences)
distinct alignments	1197	1325	1823	1950
correct	916 (77%)	1112 (85%)	1205 (66%)	1269 (66%)
usable in MT	880 (74%)	1099 (84%)	1157 (63%)	1248 (64%)

- Comparison between experiments on manually annotated treebanks and raw text

Results on course plans corpora

	PUD (100 sentences)		course plans	
	en-it	en-sv	DMI (881 sentences)	CSE (539 sentences)
distinct alignments	1197	1325	1823	1950
correct	916 (77%)	1112 (85%)	1205 (66%)	1269 (66%)
usable in MT	880 (74%)	1099 (84%)	1157 (63%)	1248 (64%)

- ❑ Comparison between experiments on manually annotated treebanks and raw text
- ❑ precision decreases, but is still higher than `fast_align`'s

Results on course plans corpora

	PUD (100 sentences)		course plans	
	en-it	en-sv	DMI (881 sentences)	CSE (539 sentences)
distinct alignments	1197	1325	1823	1950
correct	916 (77%)	1112 (85%)	1205 (66%)	1269 (66%)
usable in MT	880 (74%)	1099 (84%)	1157 (63%)	1248 (64%)

- ❑ Comparison between experiments on manually annotated treebanks and raw text
- ❑ precision decreases, but is still higher than `fast_align`'s
- ❑ recall much lower

MT experiments

- ❖ No need for an *ad hoc* grammar: extend extraction grammar with existing RGL functions

MT experiments

- ❖ No need for an *ad hoc* grammar: extend extraction grammar with existing RGL functions
- ❖ 2 bilingual lexica from course plans corpora

MT experiments

- ❖ No need for an *ad hoc* grammar: extend extraction grammar with existing RGL functions
- ❖ 2 bilingual lexica from course plans corpora
- ❖ corpus of sentences to translate generated in the GF shell
 - ❖ semi-random lexical and grammatical variations on a set of semantically plausible sentences

MT experiments

- ❖ No need for an *ad hoc* grammar: extend extraction grammar with existing RGL functions
- ❖ 2 bilingual lexica from course plans corpora
- ❖ corpus of sentences to translate generated in the GF shell
 - ❖ semi-random lexical and grammatical variations on a set of semantically plausible sentences
- ❖ metric: BLEU scores

MT experiments

- ❖ No need for an *ad hoc* grammar: extend extraction grammar with existing RGL functions
- ❖ 2 bilingual lexica from course plans corpora
- ❖ corpus of sentences to translate generated in the GF shell
 - ❖ semi-random lexical and grammatical variations on a set of semantically plausible sentences
- ❖ metric: BLEU scores
- ❖ reference translations obtained by manual postprocessing of the automatic ones

MT experiments

- ❖ No need for an *ad hoc* grammar: extend extraction grammar with existing RGL functions
- ❖ 2 bilingual lexica from course plans corpora
- ❖ corpus of sentences to translate generated in the GF shell
 - ❖ semi-random lexical and grammatical variations on a set of semantically plausible sentences
- ❖ metric: BLEU scores
- ❖ reference translations obtained by manual postprocessing of the automatic ones
 - ❖ avoid low scores due to different but equally valid lexical and grammatical choices

Results

	DMI (en-it)	CSE (en-sv)
BLEU-1 to 4	55	61
BLEU-1 to 3	63	68
BLEU-1 to 2	70	74
BLEU-1	79	81

Results

	DMI (en-it)	CSE (en-sv)
BLEU-1 to 4	55	61
BLEU-1 to 3	63	68
BLEU-1 to 2	70	74
BLEU-1	79	81

- ❖ Better results for English-Swedish (due to systematic errors in Italian)

Results

	DMI (en-it)	CSE (en-sv)
BLEU-1 to 4	55	61
BLEU-1 to 3	63	68
BLEU-1 to 2	70	74
BLEU-1	79	81

- ❖ Better results for English-Swedish (due to systematic errors in Italian)
- ❖ sentence-level scores range from 0 (sometimes due to a single semantic error) to 100

Conclusions

- ✚ Extraction technique performing consistently well even on small datasets

Conclusions

- ❖ Extraction technique performing consistently well even on small datasets
- ❖ simultaneous extraction of word, phrase, . . . alignments, incl. discontinuous expressions

Conclusions

- ❖ Extraction technique performing consistently well even on small datasets
- ❖ simultaneous extraction of word, phrase, . . . alignments, incl. discontinuous expressions
- ❖ possibility to search for specific types of correspondences, e.g. predication patterns

Conclusions

- ❖ Extraction technique performing consistently well even on small datasets
- ❖ simultaneous extraction of word, phrase, . . . alignments, incl. discontinuous expressions
- ❖ possibility to search for specific types of correspondences, e.g. predication patterns
- ❖ customizable divergence patterns

Conclusions

- ❖ Extraction technique performing consistently well even on small datasets
- ❖ simultaneous extraction of word, phrase, . . . alignments, incl. discontinuous expressions
- ❖ possibility to search for specific types of correspondences, e.g. predication patterns
- ❖ customizable divergence patterns
- ❖ output: compilable, morphology-aware GF translation lexica

Conclusions

- ❖ Extraction technique performing consistently well even on small datasets
- ❖ simultaneous extraction of word, phrase, . . . alignments, incl. discontinuous expressions
- ❖ possibility to search for specific types of correspondences, e.g. predication patterns
- ❖ customizable divergence patterns
- ❖ output: compilable, morphology-aware GF translation lexica
- ❖ require manual corrections and completions, but can significantly reduce lexicon bootstrapping time

Conclusions

- ❖ Extraction technique performing consistently well even on small datasets
- ❖ simultaneous extraction of word, phrase, . . . alignments, incl. discontinuous expressions
- ❖ possibility to search for specific types of correspondences, e.g. predication patterns
- ❖ customizable divergence patterns
- ❖ output: compilable, morphology-aware GF translation lexica
- ❖ require manual corrections and completions, but can significantly reduce lexicon bootstrapping time
- ❖ available as Haskell library + executables

Current and future work

❖ Concept Propagation

Current and future work

- ❖ Concept Propagation
 - ❖ same text in new language (equivalent to multilingual CE)

Current and future work

- ❖ Concept Propagation
 - ❖ same text in new language (equivalent to multilingual CE)
 - ❖ new text in new language (within same domain)

Current and future work

- ❖ Concept Propagation
 - ❖ same text in new language (equivalent to multilingual CE)
 - ❖ new text in new language (within same domain)
- ❖ integration with statistical tools

Current and future work

- ❖ Concept Propagation
 - ❖ same text in new language (equivalent to multilingual CE)
 - ❖ new text in new language (within same domain)
- ❖ integration with statistical tools
- ❖ postprocessing tools